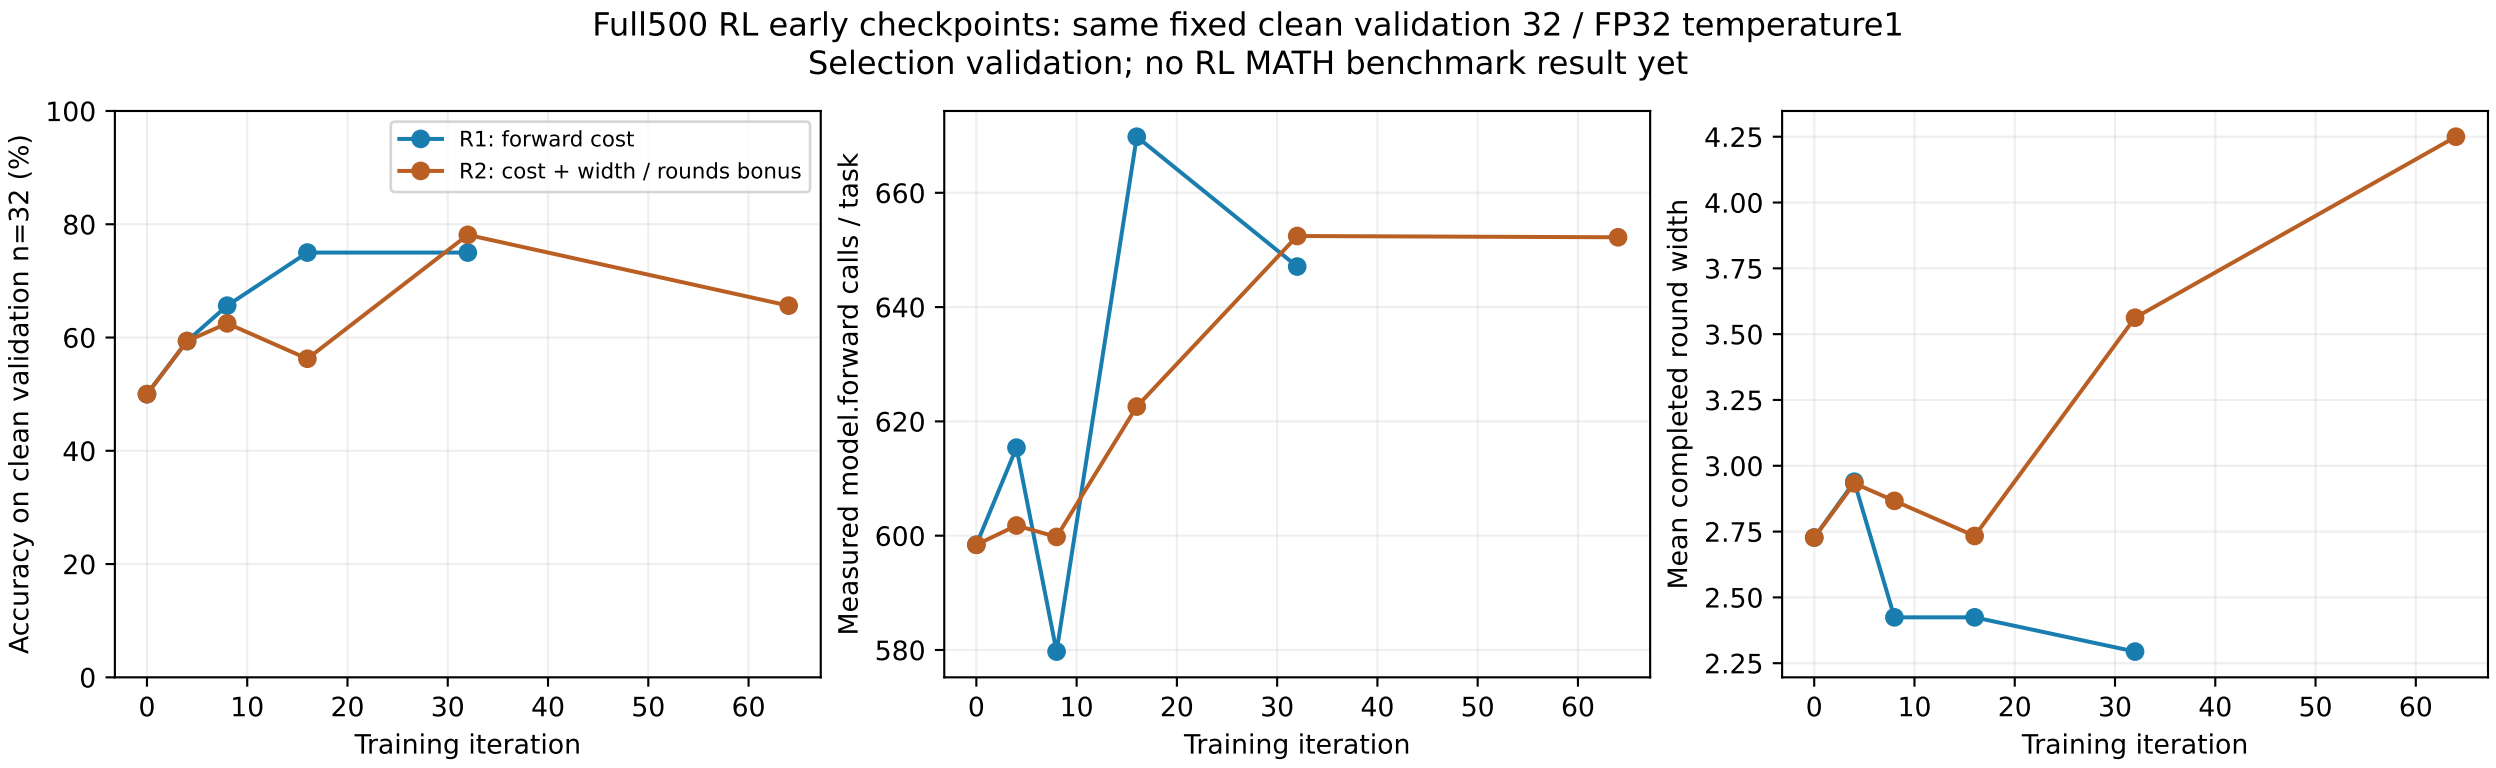
Full500 RL early checkpoints: same fixed clean validation 32 / FP32 temperature1
Selection validation; no RL MATH benchmark result yet
R1: forward cost
R2: cost + width / rounds bonus
Accuracy on clean validation n=32 (%)
Measured model.forward calls / task
Mean completed round width
Training iteration
Training iteration
Training iteration