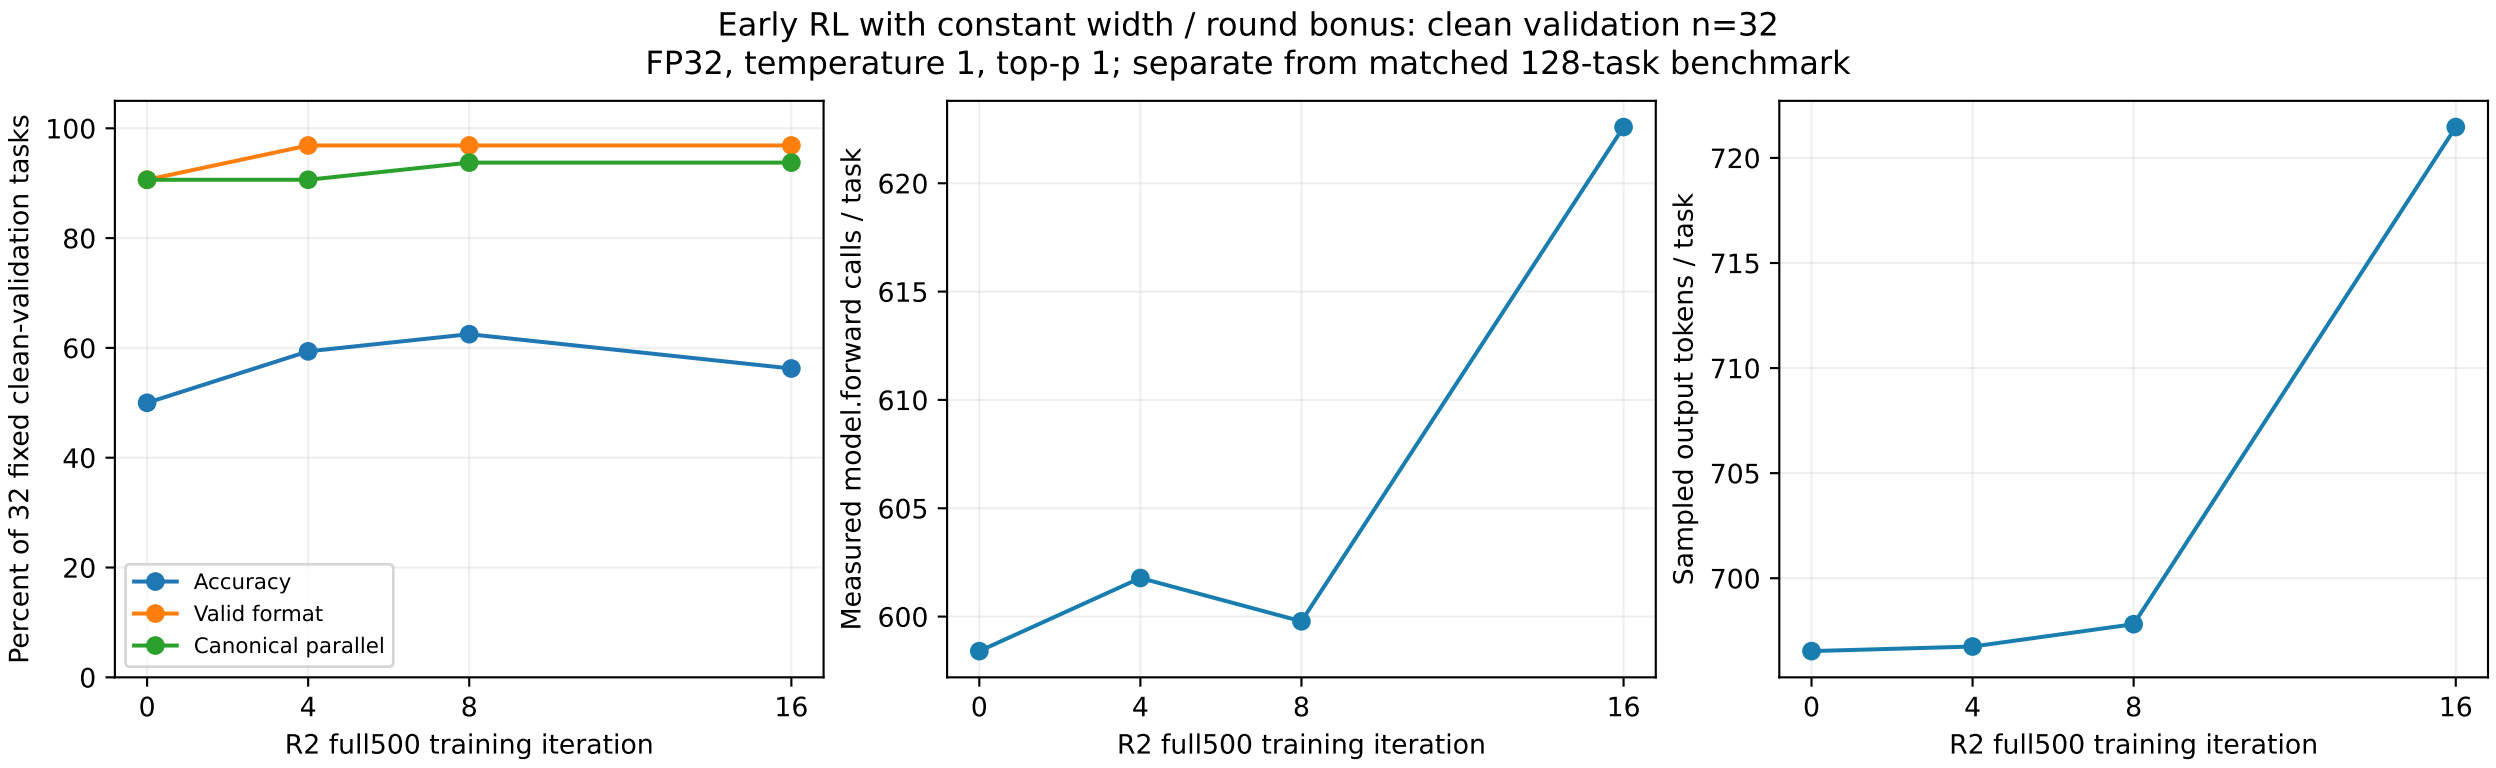
Early RL with constant width / round bonus: clean validation n=32
FP32, temperature 1, top-p 1; separate from matched 128-task benchmark
Percent of 32 fixed clean-validation tasks
Accuracy
Valid format
Canonical parallel
R2 full500 training iteration
Measured model.forward calls / task
R2 full500 training iteration
Sampled output tokens / task
R2 full500 training iteration